

```
!pip install librosa seaborn --quiet
# !pip install librosa==0.9.2
```

## ✓ 1. Load Dataset

```
import kagglehub
```

```
# Download latest version
path = kagglehub.dataset_download("uwrfkagglervavdess-emotional-speech-audio")
```

```
print("Path to dataset files:", path)
```

```
📄 Downloading from https://www.kaggle.com/api/v1/datasets/download/uwrfkagglervavdess-emotional-speech-audio?dataset\_version\_number=1.
100%|██████████| 429M/429M [00:43<00:00, 10.5MB/s]Extracting files...
```

```
Path to dataset files: /root/.cache/kagglehub/datasets/uwrfkagglervavdess-emotional-speech-audio/versions/1
```

```
import os
import sys
import warnings
import pandas as pd
import numpy as np
import librosa
import librosa.display
import seaborn as sns
import matplotlib.pyplot as plt
```

```
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.model_selection import train_test_split
```

```
from IPython.display import Audio
```

```
import keras
from keras.callbacks import ReduceLROnPlateau
from keras.models import Sequential
from keras.layers import Dense, Conv1D, MaxPooling1D, Flatten, Dropout
from keras.utils import to_categorical
```

```
warnings.filterwarnings("ignore", category=DeprecationWarning)
```

```
Ravdess = "/root/.cache/kagglehub/datasets/uwrfkagglervavdess-emotional-speech-audio/versions/1/"
```

```
# The RAVDESS filenames are structured as:
# Modality (01 = speech)
# Vocal channel (01 = speech)
# Emotion (01 = neutral, 02 = calm, 03 = happy, 04 = sad,
#          05 = angry, 06 = fearful, 07 = disgust, 08 = surprised)
# Emotional intensity (01 = normal, 02 = strong)
# Statement (01, 02)
# Repetition (01, 02)
# Actor (01 to 24)
```

```
ravdess_directory_list = os.listdir(Ravdess)
```

```
file_emotion = []
```

```
file_path = []
```

```
for dir_ in ravdess_directory_list:
```

```
    # If there are hidden files or non-directory elements, skip
```

```
    if dir_.startswith('.'):
        continue
```

```
    actor_path = os.path.join(Ravdess, dir_)
```

```
    if os.path.isdir(actor_path):
```

```
        actor_files = os.listdir(actor_path)
```

```
        for file in actor_files:
```

```
            if file.endswith('.wav'):
```

```
                part = file.split('.')[0].split('-')
```

```
                emotion_code = int(part[2])
```

```
                file_emotion.append(emotion_code)
```

```
                file_path.append(os.path.join(actor_path, file))
```

```
Ravdess_df = pd.DataFrame({'Emotions': file_emotion, 'Path': file_path})
```

```
# Instead of using inplace=True, directly assign the replaced values
```

```
Ravdess_df['Emotions'] = Ravdess_df['Emotions'].replace({
```

```
    1:'neutral', 2:'calm', 3:'happy', 4:'sad', 5:'angry',
```

```
    6:'fear', 7:'disgust', 8:'surprise'
```

```
})
```

```
# Display the first few rows to confirm
```

```
print(Ravdess_df.head())
```

```
Emotions Path
0      angry /root/.cache/kagglehub/datasets/uwrfkaggler/ra...
1    surprise /root/.cache/kagglehub/datasets/uwrfkaggler/ra...
2      happy /root/.cache/kagglehub/datasets/uwrfkaggler/ra...
3       calm /root/.cache/kagglehub/datasets/uwrfkaggler/ra...
4        sad /root/.cache/kagglehub/datasets/uwrfkaggler/ra...
```

## EDA

```
plt.figure(figsize=(10,6))
```

```
plt.title('Count of Emotions in RAVDESS', size=16)
```


```
sns.countplot(x='Emotions', data=Ravdess_df)
```

```
plt.ylabel('Count', size=12)
```

```
plt.xlabel('Emotions', size=12)
```

```
sns.despine(top=True, right=True)
```

```
plt.show()
```



```
neutral_df = Ravdess_df[Ravdess_df.Emotions == 'neutral']
```

```
num_to_add = 96
```

```
oversampled_neutral = neutral_df.sample(n=num_to_add, replace=True)
```

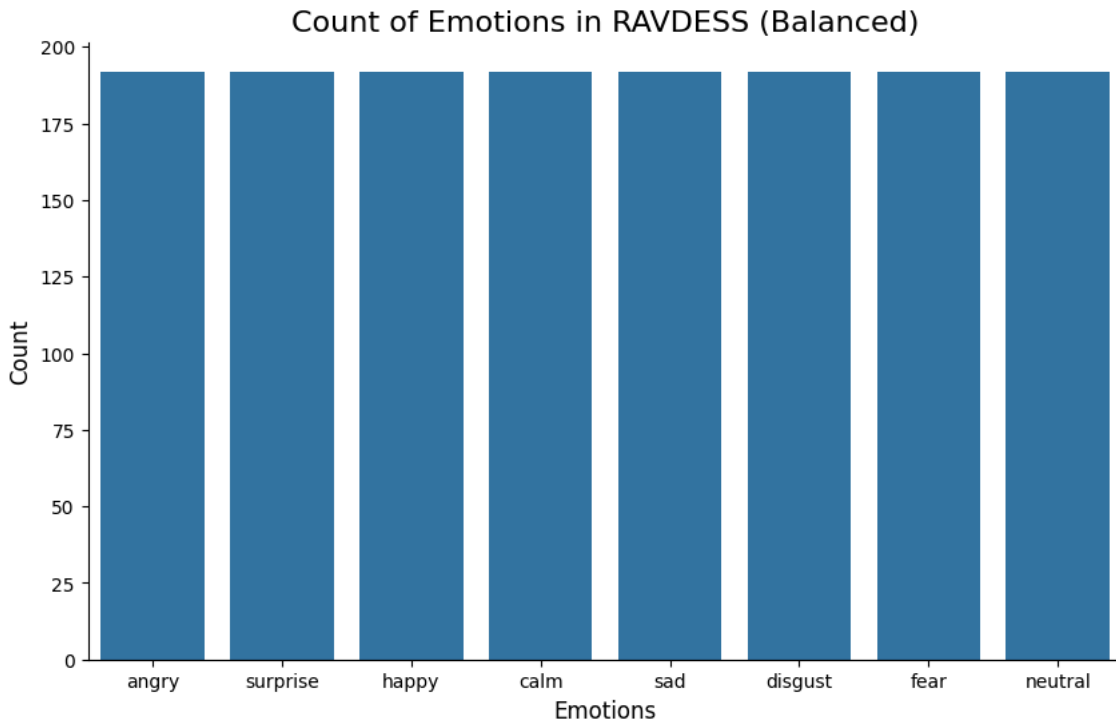
```
balanced_df = pd.concat([Ravdess_df, oversampled_neutral], ignore_index=True)
```

```
plt.figure(figsize=(10,6))
```

```
plt.title('Count of Emotions in RAVDESS (Balanced)', size=16)
```

```
sns.countplot(x='Emotions', data=balanced_df)
```

```
plt.ylabel('Count', size=12)
plt.xlabel('Emotions', size=12)
sns.despine(top=True, right=True)
plt.show()
```



```
# visualizing a few waveforms and spectrograms
```


```
def create_waveplot(data, sr, e):
    plt.figure(figsize=(10, 3))
    plt.title(f'Waveplot for audio with {e} emotion', size=15)
    librosa.display.waveshow(data, sr=sr)
    plt.show()

def create_spectrogram(data, sr, e):
    X = librosa.stft(data)
    Xdb = librosa.amplitude_to_db(abs(X))
    plt.figure(figsize=(12, 3))
    plt.title(f'Spectrogram for audio with {e} emotion', size=15)
    librosa.display.specshow(Xdb, sr=sr, x_axis='time', y_axis='hz')
    plt.colorbar()
    plt.show()
```


```
# Example visualization
example_emotion = 'happy'
example_path = Ravdess_df[Ravdess_df.Emotions == example_emotion].iloc[0].Path
data, sampling_rate = librosa.load(example_path)
create_waveplot(data, sampling_rate, example_emotion)
create_spectrogram(data, sampling_rate, example_emotion)
Audio(example_path)
```



Waveplot for audio with happy emotion



Spectrogram for audio with happy emotion



0:00 / 0:03

### 3. Data Augmentation Functions

```
# Data Augmentation Functions
def noise(data):
    noise_amp = 0.035 * np.random.uniform() * np.amax(data)
    data = data + noise_amp * np.random.normal(size=data.shape[0])
    return data

def stretch(data, rate=0.8):
    # Directly apply time stretching on the waveform with keyword arguments
    # This ensures that the correct arguments are passed even if there's a namespace conflict
    y_stretched = librosa.effects.time_stretch(y=data, rate=rate)
    return y_stretched

def pitch(data, sr, pitch_factor=0.7):
    return librosa.effects.pitch_shift(y=data, sr=sr, n_steps=pitch_factor)
```

### 4. Feature Extraction

```
def extract_features(data, sr):
    result = np.array([])

    # ZCR
    zcr = np.mean(librosa.feature.zero_crossing_rate(y=data).T, axis=0)
    result = np.hstack((result, zcr))

    # Chroma
    stft = np.abs(librosa.stft(data))
    chroma_stft = np.mean(librosa.feature.chroma_stft(S=stft, sr=sr).T, axis=0)
    result = np.hstack((result, chroma_stft))

    # MFCC
    mfcc = np.mean(librosa.feature.mfcc(y=data, sr=sr).T, axis=0)
    result = np.hstack((result, mfcc))

    # RMS
    rms = np.mean(librosa.feature.rms(y=data).T, axis=0)
    result = np.hstack((result, rms))

    # MelSpectrogram
    mel = np.mean(librosa.feature.melspectrogram(y=data, sr=sr).T, axis=0)
    result = np.hstack((result, mel))

    return result
```

```

def get_features(path):
    data, sr = librosa.load(path, duration=2.5, offset=0.6)
    # original
    res1 = extract_features(data, sr)
    result = np.array(res1)

    # with noise
    noise_data = noise(data)
    res2 = extract_features(noise_data, sr)
    result = np.vstack((result, res2))

    # with stretching + pitching
    stretched_data = stretch(data)
    pitched_data = pitch(stretched_data, sr)
    res3 = extract_features(pitched_data, sr)
    result = np.vstack((result, res3))

    return result

X, Y = [], []
for path, emotion in zip(Ravdess_df.Path, Ravdess_df.Emotions):
    feature = get_features(path)
    for ele in feature:
        X.append(ele)
        Y.append(emotion)

X = np.array(X)
Y = np.array(Y)

# One-hot encode targets
encoder = OneHotEncoder()
Y = encoder.fit_transform(Y.reshape(-1,1)).toarray()

# Train-test split
x_train, x_test, y_train, y_test = train_test_split(X, Y, random_state=0, shuffle=True)

# Normalize data
scaler = StandardScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.transform(x_test)

# Add channel dimension
x_train = np.expand_dims(x_train, axis=2)
x_test = np.expand_dims(x_test, axis=2)

# 5. Model Building
model = Sequential()
model.add(Conv1D(256, kernel_size=5, strides=1, padding='same', activation='relu', input_shape=(x_train.shape[1], 1)))
model.add(MaxPooling1D(pool_size=5, strides=2, padding='same'))

model.add(Conv1D(256, kernel_size=5, strides=1, padding='same', activation='relu'))
model.add(MaxPooling1D(pool_size=5, strides=2, padding='same'))

model.add(Conv1D(128, kernel_size=5, strides=1, padding='same', activation='relu'))
model.add(MaxPooling1D(pool_size=5, strides=2, padding='same'))
model.add(Dropout(0.2))

model.add(Conv1D(64, kernel_size=5, strides=1, padding='same', activation='relu'))
model.add(MaxPooling1D(pool_size=5, strides=2, padding='same'))

model.add(Flatten())
model.add(Dense(units=32, activation='relu'))
model.add(Dropout(0.3))
model.add(Dense(units=8, activation='softmax')) # 8 emotions in RAVDESS

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
model.summary()

```

```

/usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/'i
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Model: "sequential"

```

Layer (type)	Output Shape	Param #
conv1d (Conv1D)	(None, 162, 256)	1,536
max_pooling1d (MaxPooling1D)	(None, 81, 256)	0
conv1d_1 (Conv1D)	(None, 81, 256)	327,936
max_pooling1d_1 (MaxPooling1D)	(None, 41, 256)	0
conv1d_2 (Conv1D)	(None, 41, 128)	163,968
max_pooling1d_2 (MaxPooling1D)	(None, 21, 128)	0
dropout (Dropout)	(None, 21, 128)	0
conv1d_3 (Conv1D)	(None, 21, 64)	41,024
max_pooling1d_3 (MaxPooling1D)	(None, 11, 64)	0
flatten (Flatten)	(None, 704)	0
dense (Dense)	(None, 32)	22,560
dropout_1 (Dropout)	(None, 32)	0
dense_1 (Dense)	(None, 8)	264

```

Total params: 557,288 (2.13 MB)
Trainable params: 557,288 (2.13 MB)
Non-trainable params: 0 (0.00 B)

```

```

# Reduce learning rate on plateau
rlrp = ReduceLROnPlateau(monitor='loss', factor=0.4, patience=2, min_lr=1e-7)

history = model.fit(x_train, y_train, batch_size=64, epochs=50, validation_data=(x_test, y_test), callbacks=[rlrp])

```

```

Epoch 1/50
51/51 ----- 17s 172ms/step - accuracy: 0.1867 - loss: 2.0301 - val_accuracy: 0.2139 - val_loss: 1.9193 - learning_rate: 0.001
Epoch 2/50
51/51 ----- 6s 11ms/step - accuracy: 0.2274 - loss: 1.9485 - val_accuracy: 0.2481 - val_loss: 1.8751 - learning_rate: 0.001
Epoch 3/50
51/51 ----- 0s 8ms/step - accuracy: 0.2470 - loss: 1.8844 - val_accuracy: 0.3056 - val_loss: 1.7819 - learning_rate: 0.001
Epoch 4/50
51/51 ----- 1s 8ms/step - accuracy: 0.2934 - loss: 1.8028 - val_accuracy: 0.3250 - val_loss: 1.7752 - learning_rate: 0.001
Epoch 5/50
51/51 ----- 0s 8ms/step - accuracy: 0.2748 - loss: 1.8130 - val_accuracy: 0.3231 - val_loss: 1.7074 - learning_rate: 0.001
Epoch 6/50
51/51 ----- 1s 9ms/step - accuracy: 0.3148 - loss: 1.7511 - val_accuracy: 0.3583 - val_loss: 1.6632 - learning_rate: 0.001
Epoch 7/50
51/51 ----- 0s 9ms/step - accuracy: 0.3357 - loss: 1.6877 - val_accuracy: 0.3870 - val_loss: 1.6038 - learning_rate: 0.001
Epoch 8/50
51/51 ----- 1s 9ms/step - accuracy: 0.3568 - loss: 1.6335 - val_accuracy: 0.3694 - val_loss: 1.5926 - learning_rate: 0.001
Epoch 9/50
51/51 ----- 1s 9ms/step - accuracy: 0.3794 - loss: 1.6252 - val_accuracy: 0.4046 - val_loss: 1.5690 - learning_rate: 0.001
Epoch 10/50
51/51 ----- 0s 8ms/step - accuracy: 0.3854 - loss: 1.5927 - val_accuracy: 0.4000 - val_loss: 1.5185 - learning_rate: 0.001
Epoch 11/50
51/51 ----- 0s 8ms/step - accuracy: 0.4082 - loss: 1.5451 - val_accuracy: 0.4056 - val_loss: 1.5466 - learning_rate: 0.001
Epoch 12/50
51/51 ----- 0s 8ms/step - accuracy: 0.4229 - loss: 1.5106 - val_accuracy: 0.4056 - val_loss: 1.5417 - learning_rate: 0.001
Epoch 13/50
51/51 ----- 0s 8ms/step - accuracy: 0.4265 - loss: 1.4831 - val_accuracy: 0.4444 - val_loss: 1.4686 - learning_rate: 0.001
Epoch 14/50
51/51 ----- 1s 8ms/step - accuracy: 0.4406 - loss: 1.4399 - val_accuracy: 0.4380 - val_loss: 1.4510 - learning_rate: 0.001
Epoch 15/50
51/51 ----- 0s 8ms/step - accuracy: 0.4339 - loss: 1.4576 - val_accuracy: 0.4657 - val_loss: 1.4081 - learning_rate: 0.001
Epoch 16/50
51/51 ----- 0s 8ms/step - accuracy: 0.4768 - loss: 1.3825 - val_accuracy: 0.4833 - val_loss: 1.3879 - learning_rate: 0.001
Epoch 17/50
51/51 ----- 0s 8ms/step - accuracy: 0.4862 - loss: 1.3540 - val_accuracy: 0.4731 - val_loss: 1.3616 - learning_rate: 0.001
Epoch 18/50
51/51 ----- 1s 8ms/step - accuracy: 0.4840 - loss: 1.3347 - val_accuracy: 0.4954 - val_loss: 1.3376 - learning_rate: 0.001
Epoch 19/50
51/51 ----- 1s 8ms/step - accuracy: 0.5100 - loss: 1.2925 - val_accuracy: 0.4935 - val_loss: 1.3271 - learning_rate: 0.001
Epoch 20/50
51/51 ----- 0s 8ms/step - accuracy: 0.5163 - loss: 1.2626 - val_accuracy: 0.4509 - val_loss: 1.4760 - learning_rate: 0.001
Epoch 21/50
51/51 ----- 0s 8ms/step - accuracy: 0.5075 - loss: 1.3011 - val_accuracy: 0.5194 - val_loss: 1.2879 - learning_rate: 0.001
Epoch 22/50
51/51 ----- 1s 8ms/step - accuracy: 0.5488 - loss: 1.1943 - val_accuracy: 0.5287 - val_loss: 1.2709 - learning_rate: 0.001
Epoch 23/50
51/51 ----- 0s 8ms/step - accuracy: 0.5647 - loss: 1.1645 - val_accuracy: 0.5343 - val_loss: 1.2534 - learning_rate: 0.001
Epoch 24/50
51/51 ----- 1s 8ms/step - accuracy: 0.5666 - loss: 1.1170 - val_accuracy: 0.5454 - val_loss: 1.2320 - learning_rate: 0.001
Epoch 25/50
51/51 ----- 0s 8ms/step - accuracy: 0.6124 - loss: 1.0368 - val_accuracy: 0.5296 - val_loss: 1.3365 - learning_rate: 0.001
Epoch 26/50
51/51 ----- 0s 8ms/step - accuracy: 0.5845 - loss: 1.1098 - val_accuracy: 0.5676 - val_loss: 1.1863 - learning_rate: 0.001
Epoch 27/50

```

```

51/51 ————— 0s 8ms/step - accuracy: 0.6299 - loss: 0.9866 - val_accuracy: 0.5111 - val_loss: 1.3045 - learning_rate: 0
Epoch 28/50
51/51 ————— 1s 8ms/step - accuracy: 0.6123 - loss: 1.0166 - val_accuracy: 0.5324 - val_loss: 1.2873 - learning_rate: 0
Epoch 29/50
51/51 ————— 0s 8ms/step - accuracy: 0.6288 - loss: 0.9440 - val_accuracy: 0.5565 - val_loss: 1.2102 - learning_rate: 0

```

# 6. Evaluation

```

test_loss, test_acc = model.evaluate(x_test, y_test)
print(f"Test Accuracy: {test_acc*100:.2f}%")

```

# Plotting Accuracy and Loss

```

epochs = range(1,51)
train_acc = history.history['accuracy']
train_loss = history.history['loss']
val_acc = history.history['val_accuracy']
val_loss = history.history['val_loss']

```

```

plt.figure(figsize=(14,5))
plt.subplot(1,2,1)
plt.plot(epochs, train_loss, 'b', label='Training Loss')
plt.plot(epochs, val_loss, 'r', label='Testing Loss')
plt.title('Training & Testing Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

```

```


plt.subplot(1,2,2)
plt.plot(epochs, train_acc, 'b', label='Training Accuracy')
plt.plot(epochs, val_acc, 'r', label='Testing Accuracy')
plt.title('Training & Testing Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

```

```

34/34 ————— 2s 25ms/step - accuracy: 0.6154 - loss: 1.2844
Test Accuracy: 60.37%

```



# Predictions and Confusion Matrix

```

pred_test = model.predict(x_test)
y_pred = encoder.inverse_transform(pred_test)
y_true = encoder.inverse_transform(y_test)

```

```


34/34 ————— 1s 16ms/step

```

```

cm = confusion_matrix(y_true, y_pred)
plt.figure(figsize=(12, 10))
cm_df = pd.DataFrame(cm, index=encoder.categories_[0], columns=encoder.categories_[0])
sns.heatmap(cm_df, annot=True, cmap='Blues', fmt='g')
plt.title('Confusion Matrix')
plt.xlabel('Predicted Labels')
plt.ylabel('Actual Labels')
plt.show()

```



```
print(classification_report(y_true, y_pred))
```

↻

	precision	recall	f1-score	support
angry	0.63	0.72	0.67	149
calm	0.68	0.78	0.72	139
disgust	0.52	0.60	0.56	149
fear	0.60	0.71	0.65	153
happy	0.58	0.41	0.48	145
neutral	0.48	0.53	0.50	74
sad	0.63	0.50	0.56	137
surprise	0.69	0.53	0.60	134
accuracy			0.60	1080
macro avg	0.60	0.60	0.59	1080
weighted avg	0.61	0.60	0.60	1080

```
import joblib
from keras.models import load_model
from google.colab import files
```

```
model.save('emotion_model.keras')
joblib.dump(scaler, 'scaler.pkl')
joblib.dump(encoder, 'encoder.pkl')
```

↻ ['encoder.pkl']

```
files.download('emotion_model.keras')
files.download('scaler.pkl')
files.download('encoder.pkl')
```

↻

```
#Trying to improve
```

```
Tess = kagglehub.dataset_download("ejlok1/toronto-emotional-speech-set-tess")
TESS_directory_list = os.listdir(Tess)
Crema = kagglehub.dataset_download("ejlok1/cremad")
```



```
CremaD_directory_list = os.listdir(Crema)
Savee = kagglehub.dataset_download("ejlok1/surrey-audiovisual-expressed-emotion-savee")
Savee_directory_list = os.listdir(Savee)
```

```
Downloading from https://www.kaggle.com/api/v1/datasets/download/ejlok1/toronto-emotional-speech-set-tess?dataset\_version\_number=1...
100%|██████████| 428M/428M [00:08<00:00, 51.4MB/s]Extracting files...

Downloading from https://www.kaggle.com/api/v1/datasets/download/ejlok1/cremad?dataset\_version\_number=1...
100%|██████████| 451M/451M [00:03<00:00, 149MB/s]Extracting files...

Downloading from https://www.kaggle.com/api/v1/datasets/download/ejlok1/surrey-audiovisual-expressed-emotion-savee?dataset\_version\_nu
100%|██████████| 107M/107M [00:01<00:00, 107MB/s] Extracting files...
```

```
# Crema = "/path_to_crema/"
# Tess = "/path_to_tess/"
# Savee = "/path_to_savee/"
```

```
# Load CREMA DataFrame (Adjust paths accordingly)
crema_emotion = []
crema_path = []
for file in os.listdir(Crema):
    if file.endswith('.wav'):
        file_path = os.path.join(Crema, file)
        part = file.split('_')
        # Map according to the CREMA naming scheme
        if part[2] == 'SAD':
            crema_emotion.append('sad')
        elif part[2] == 'ANG':
            crema_emotion.append('angry')
        elif part[2] == 'DIS':
            crema_emotion.append('disgust')
        elif part[2] == 'FEA':
            crema_emotion.append('fear')
        elif part[2] == 'HAP':
            crema_emotion.append('happy')
        elif part[2] == 'NEU':
            crema_emotion.append('neutral')
        else:
            crema_emotion.append('Unknown')
        crema_path.append(file_path)
Crema_df = pd.DataFrame({'Emotions': crema_emotion, 'Path': crema_path})
```


```
# Load TESS DataFrame
tess_emotion = []
tess_path = []
for folder in os.listdir(Tess):
    folder_path = os.path.join(Tess, folder)
    if os.path.isdir(folder_path):
        for file in os.listdir(folder_path):
            if file.endswith('.wav'):
                part = file.split('.')[0].split('_')[-1]
                if part == 'ps':
                    tess_emotion.append('surprise')
                else:
                    tess_emotion.append(part)
                tess_path.append(os.path.join(folder_path, file))
Tess_df = pd.DataFrame({'Emotions': tess_emotion, 'Path': tess_path})
```

```
# Load SAVEE DataFrame
savee_emotion = []
savee_path = []
for file in os.listdir(Savee):
    if file.endswith('.wav'):
        file_path = os.path.join(Savee, file)
        part = file.split('_')[1][:6]
        if part == 'a':
            savee_emotion.append('angry')
        elif part == 'd':
            savee_emotion.append('disgust')
        elif part == 'f':
            savee_emotion.append('fear')
        elif part == 'h':
            savee_emotion.append('happy')
        elif part == 'n':
            savee_emotion.append('neutral')
        elif part == 'sa':
            savee_emotion.append('sad')
        else:
            savee_emotion.append('surprise')
        savee_path.append(file_path)
Savee_df = pd.DataFrame({'Emotions': savee_emotion, 'Path': savee_path})
```

```
# Combine all DataFrames
data_path = pd.concat([Ravdess_df, Crema_df, Tess_df, Savee_df], ignore_index=True)
```

```
# Perform EDA
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(12,7))
sns.countplot(x='Emotions', data=data_path)
plt.title("Combined Emotion Counts")
plt.xticks(rotation=45)
plt.show()
```




```
neutral_df = Ravdess_df[Ravdess_df.Emotions == 'neutral']
num_to_add = 96
oversampled_neutral = neutral_df.sample(n=num_to_add, replace=True)
balanced_df = pd.concat([Ravdess_df, oversampled_neutral], ignore_index=True)
```

```
plt.figure(figsize=(10,6))
plt.title('Count of Emotions in Dataset (Balanced)', size=16)
sns.countplot(x='Emotions', data=balanced_df)
plt.ylabel('Count', size=12)
plt.xlabel('Emotions', size=12)
sns.despine(top=True, right=True)
plt.show()
```



## Count of Emotions in RAVDESS (Balanced)



```
# Feature extraction (assuming get_features is defined as before)
X, Y = [], []
for path, emotion in zip(data_path.Path, data_path.Emotions):
    feats = get_features(path) # uses the same augmentations as before
    for ele in feats:
        X.append(ele)
        Y.append(emotion)

X = np.array(X)
Y = np.array(Y)

# One-hot encode targets
encoder = OneHotEncoder()
Y = encoder.fit_transform(Y.reshape(-1,1)).toarray()

# Train-test split
x_train, x_test, y_train, y_test = train_test_split(X, Y, random_state=0, shuffle=True)

# Normalize data
scaler = StandardScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.transform(x_test)

# Add channel dimension
x_train = np.expand_dims(x_train, axis=2)
x_test = np.expand_dims(x_test, axis=2)

# 5. Model Building
model = Sequential()
model.add(Conv1D(256, kernel_size=5, strides=1, padding='same', activation='relu', input_shape=(x_train.shape[1], 1)))
model.add(MaxPooling1D(pool_size=5, strides=2, padding='same'))

model.add(Conv1D(256, kernel_size=5, strides=1, padding='same', activation='relu'))
model.add(MaxPooling1D(pool_size=5, strides=2, padding='same'))

model.add(Conv1D(128, kernel_size=5, strides=1, padding='same', activation='relu'))
model.add(MaxPooling1D(pool_size=5, strides=2, padding='same'))
model.add(Dropout(0.2))

model.add(Conv1D(64, kernel_size=5, strides=1, padding='same', activation='relu'))
model.add(MaxPooling1D(pool_size=5, strides=2, padding='same'))

model.add(Flatten())
model.add(Dense(units=32, activation='relu'))
model.add(Dropout(0.3))
model.add(Dense(units=8, activation='softmax')) # 8 emotions in RAVDESS

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
model.summary()
```

```

/usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/'i
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Model: "sequential_2"

```

Layer (type)	Output Shape	Param #
conv1d_8 (Conv1D)	(None, 162, 256)	1,536
max_pooling1d_8 (MaxPooling1D)	(None, 81, 256)	0
conv1d_9 (Conv1D)	(None, 81, 256)	327,936
max_pooling1d_9 (MaxPooling1D)	(None, 41, 256)	0
conv1d_10 (Conv1D)	(None, 41, 128)	163,968
max_pooling1d_10 (MaxPooling1D)	(None, 21, 128)	0
dropout_4 (Dropout)	(None, 21, 128)	0
conv1d_11 (Conv1D)	(None, 21, 64)	41,024
max_pooling1d_11 (MaxPooling1D)	(None, 11, 64)	0
flatten_2 (Flatten)	(None, 704)	0
dense_4 (Dense)	(None, 32)	22,560
dropout_5 (Dropout)	(None, 32)	0
dense_5 (Dense)	(None, 8)	264

```

Total params: 557,288 (2.13 MB)
Trainable params: 557,288 (2.13 MB)
Non-trainable params: 0 (0.00 B)

```

```

# Reduce learning rate on plateau
rlrp = ReduceLRonPlateau(monitor='loss', factor=0.4, patience=2, min_lr=1e-7)

history = model.fit(x_train, y_train, batch_size=64, epochs=50, validation_data=(x_test, y_test), callbacks=[rlrp])

```

```

Epoch 22/50
51/51 ----- 0s 8ms/step - accuracy: 0.5552 - loss: 1.1789 - val_accuracy: 0.5148 - val_loss: 1.3089 - learning_rate: 0
Epoch 23/50
51/51 ----- 0s 8ms/step - accuracy: 0.5481 - loss: 1.1554 - val_accuracy: 0.5185 - val_loss: 1.2557 - learning_rate: 0
Epoch 24/50
51/51 ----- 1s 8ms/step - accuracy: 0.5853 - loss: 1.0713 - val_accuracy: 0.5176 - val_loss: 1.3189 - learning_rate: 0
Epoch 25/50
51/51 ----- 0s 8ms/step - accuracy: 0.5997 - loss: 1.0485 - val_accuracy: 0.5417 - val_loss: 1.2949 - learning_rate: 0
Epoch 26/50
51/51 ----- 0s 8ms/step - accuracy: 0.6101 - loss: 1.0114 - val_accuracy: 0.5704 - val_loss: 1.2342 - learning_rate: 0
Epoch 27/50
51/51 ----- 0s 8ms/step - accuracy: 0.6291 - loss: 0.9632 - val_accuracy: 0.5481 - val_loss: 1.2976 - learning_rate: 0

```

Epoch 49/50  
51/51 — 0s 8ms/step - accuracy: 0.8773 - loss: 0.3293 - val\_accuracy: 0.6389 - val\_loss: 1.4464 - learning\_rate: 4e-05  
Epoch 50/50  
51/51 — 1s 8ms/step - accuracy: 0.8955 - loss: 0.3116 - val\_accuracy: 0.6472 - val\_loss: 1.4701 - learning\_rate: 4e-05

## # 6. Evaluation

```
test_loss, test_acc = model.evaluate(x_test, y_test)
print(f"Test Accuracy: {test_acc*100:.2f}%")
```


## # Plotting Accuracy and Loss

```
epochs = range(1,51)
train_acc = history.history['accuracy']
train_loss = history.history['loss']
val_acc = history.history['val_accuracy']
val_loss = history.history['val_loss']

plt.figure(figsize=(14,5))
plt.subplot(1,2,1)
plt.plot(epochs, train_loss, 'b', label='Training Loss')
plt.plot(epochs, val_loss, 'r', label='Testing Loss')
plt.title('Training & Testing Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
```

```
plt.subplot(1,2,2)
plt.plot(epochs, train_acc, 'b', label='Training Accuracy')
plt.plot(epochs, val_acc, 'r', label='Testing Accuracy')
plt.title('Training & Testing Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```

34/34 — 1s 13ms/step - accuracy: 0.6458 - loss: 1.4286  
Test Accuracy: 64.72%



## # Predictions and Confusion Matrix


```
pred_test = model.predict(x_test)
y_pred = encoder.inverse_transform(pred_test)
y_true = encoder.inverse_transform(y_test)
```

34/34 — 1s 9ms/step

```
cm = confusion_matrix(y_true, y_pred)
plt.figure(figsize=(12, 10))
cm_df = pd.DataFrame(cm, index=encoder.categories_[0], columns=encoder.categories_[0])
sns.heatmap(cm_df, annot=True, cmap='Blues', fmt='g')
plt.title('Confusion Matrix')
plt.xlabel('Predicted Labels')
plt.ylabel('Actual Labels')
plt.show()
print(classification_report(y_true, y_pred))
```



Confusion Matrix



	precision	recall	f1-score	support
angry	0.76	0.67	0.71	149
calm	0.73	0.81	0.76	139
disgust	0.68	0.71	0.69	149
fear	0.65	0.66	0.66	153
happy	0.57	0.58	0.57	145
neutral	0.43	0.53	0.48	74
sad	0.59	0.63	0.61	137
surprise	0.72	0.53	0.61	134
accuracy			0.65	1080
macro avg	0.64	0.64	0.64	1080
weighted avg	0.65	0.65	0.65	1080